Microscopy Cell Nuclei Segmentation with U-Nets

Humza Iqbal

December 19th, 2023

1 Introduction to Problem

In microscopy, cell segmentation poses a fundamental challenge in image analysis, as it involves the precise identification and delineation of individual cells within complex biological structures. Accurate cell segmentation is crucial for various biomedical applications, such as understanding cell behavior, tracking cellular processes, and diagnosing diseases.

There are many difficult problems when processing these images. For example, cells may have compartments like nuclei, which we may not want to reclassify. The position and orientation of cells within images may vary, making it challenging to design handcrafted features that are invariant to these transformations. These problems are solved with Convolution Neural Networks, which have the necessary functional capability and expressivity.

To the main point, The central challenge in cell segmentation often revolves around the scarcity of a substantial number of samples. However, U-nets offer a promising solution to address this limitation. U-nets play a crucial role by implementing augmentations on the input images, effectively creating a pseudo-bootstrap approach. This pseudo-bootstrap involves applying various transformations to the input images, introducing diversity and variance. By doing so, U-nets generate an augmented dataset that significantly expands the available training data.

The augmentation process enhances the robustness of the U-net model by exposing it to a broader range of scenarios and variations in the input data. This not only mitigates the challenge of limited samples but also allows the network to learn more effectively from the augmented dataset. The augmentation-driven pseudo-bootstrap serves as a practical workaround when acquiring an extensive set of original samples is impractical.

Moreover, the effectiveness of U-nets in cell segmentation extends beyond their role in data augmentation. Olaf Ronneberger's comprehensive article, "U-Net: Convolutional Networks for Biomedical Image Segmentation," delves into various aspects of U-net performance. Ronneberger provides detailed insights into how U-nets leverage their unique architecture to excel in biomedical image segmentation tasks.

2   Technical Details and Formal Description

A Convolutional Neural Network (CNN) is a specialized type of artificial neural network designed for tasks such as image recognition and processing. It excels at automatically learning hierarchical feature representations from input data.

Convolutional Layers utilize convolution operations, where small filters (kernels) slide across the input data, computing dot products at each step. This enables the network to learn local patterns and features. The linearity and commutativity of the convolution operator play crucial roles in this process, ensuring that the order of operations does not affect the final result. Mathematically, the convolution operator can be represented as,

$$(f * g)(t) = \int f(\tau)g(t - \tau)d\tau$$

where f represents the input data, g represents the filter, and t represents the position.

Applied after convolution operations, activation functions (e.g., ReLU) introduce non-linearity, enabling the network to learn complex relationships. This nonlinearity is essential for capturing intricate patterns and variations in the data, especially since the convolution operator is linear. The ReLU is defined below,

$$\text{ReLU}(x) = \max(0, x)$$

Pooling layers downsample spatial dimensions, reducing the number of parameters and computational load. Techniques like max-pooling retain the maximum value in local regions, contributing to translation invariance by focusing on the most significant features in different input parts.

$$\text{MaxPooling}(x, y) = \max_{i \in N(x,y)} f(i)$$

Where $N(x, y)$ is in the set of indices in the local region (x, y).

CNN's main goal is to promote properties of invariance and equivariance. CNNs exhibit translation invariance, meaning they can recognize patterns regardless of their position in the input.

$$(f * g)(t) = (f * g)(t - t_0)$$

Equivariance is an extension, suggesting that the network's output varies predictably with input transformations. Suppose f and g are related by transformation T. Then we know, that $(f * g)(t)$ and $(Tf * g)(t)$ are related. This property is crucial for tasks where the spatial relationships between features are essential.

Typically found in the later stages of a CNN, fully connected layers connect every neuron, enabling the network to make final predictions. We have then the following objective,

$$F(x) = W_L \cdot \text{ReLU}(W_{L-1} \cdot \ldots \cdot \text{ReLU}(W_1 \cdot x + b_1) + b_{L-1}) + b_L$$

3   Results, Limitations & Improvements

4   Implementation and Reproduction

References

Hastie, T., Qian, J., & Tay, K. (2023, March 27). *An introduction to 'glmnet'.* glmnet
https://glmnet.stanford.edu/articles/glmnet.html